

Propositional Probability II: Reasoning¹

Steven Holtzen

s.holtzen@northeastern.edu

September 15, 2023

¹ CS7470 Fall 2023: Foundations of Probabilistic Programming.

Now we are ready to ask: *can we use propositional logic as the basis for an effective probabilistic model?* The first question is: what is our sample space? It's clear so far that a logical choice of sample space is the set of all possibly instances for a fixed set of Boolean formulae; this sample space will have size 2^n , where n is the number of propositional variables. Now, if we want to efficiently evaluate queries over this sample space, we need two more components:

1. A way to efficiently represent a probability distribution on the sample space;
2. A way to efficiently represent queries over that sample space.

Definition 1 (Fully factorized probabilistic model). *Let X_1, X_2, \dots, X_n be jointly independent random variables (i.e., for any pair X_i, X_j , it is the case that $X_i \perp\!\!\!\perp X_j$). Then, a fully-factorized probabilistic model is a collection of n probability lookup tables $\Pr(X_i)$, one for each i . The joint probability is computed as $\Pr(X_1, X_2, \dots, X_n) \triangleq \prod_i^n \Pr(X_i)$.*

1 A Propositional PPL (PROP)

Syntax of PROP:

```
1  $\varphi ::= x \mid \varphi \wedge \varphi \mid \varphi \vee \varphi \mid \neg \varphi \mid \text{true} \mid \text{false}$ 
2  $w ::= [x \mapsto \theta_x, y \mapsto \theta_y, \dots]$ 
3  $p ::= (\varphi, w)$ 
```

Example PROP program:

```
1  $(x \vee y, [x \mapsto 0.1, y \mapsto 0.2])$ 
```

We will assume that all well-formed PROP programs give a probability to each variable, and in those note we will restrict ourselves to considering well-formed programs.

- Denotational semantics of PROP: we want to associate every program in PROP with the probability that φ holds according to the fully-factorized distribution described by w
- **Goal:** Define a map $\llbracket \cdot \rrbracket : p \rightarrow [0, 1]$ that is a map from PROP terms to real values.
- We will need semantics for φ and the map m in order to interpret p . They have the following types:

For more on the history and context of different styles of semantics, see Pierce [2002, Chapter 3]

- $\llbracket \varphi \rrbracket$ maps propositional formulae to the set of all instances that model them
- $\llbracket w \rrbracket$ produces a map from literals (propositional variables and a truth assignment) to real values

$$\llbracket \varphi \rrbracket \triangleq \{I \mid I \models \varphi\} \quad (1)$$

$$\llbracket [x \mapsto \theta, \dots] \rrbracket(x) \triangleq \theta \quad (2)$$

$$\llbracket [x \mapsto \theta, \dots] \rrbracket(\bar{x}) \triangleq 1 - \theta \quad (3)$$

$$\llbracket [] \rrbracket(x) \triangleq 0 \quad (4)$$

$$\llbracket [y \mapsto \theta_y, r] \rrbracket(x) \triangleq \llbracket [r] \rrbracket(x) \quad (5)$$

We need to assume a propositional universe so that these equations are well-defined. Assume that all instances are defined on all free variables in p .

- We define the probability of an instance $\llbracket w \rrbracket(I)$ as the product of probabilities of each variable in the instance. For instance,

$$\llbracket (x \mapsto 0.1, y \mapsto 0.3) \rrbracket(x, \bar{y}) = 0.1 * 0.7.$$

Note that we are interpreting the probability of a negated variable as 1 minus the probability of the positive variable.

- Now we are ready to interpret PROP programs as the sum of the probabilities of each model:

$$\llbracket (\varphi, w) \rrbracket \triangleq \sum_{I \in \llbracket \varphi \rrbracket} \prod_{\ell \in I} \llbracket w \rrbracket(\ell). \quad (6)$$

- Is this a *useful* PPL? What kinds of programs can we write in it?

2 Big-step semantics

- Our denotational model does not tell us how to efficiently compute probabilities
- **Goal:** given a PROP program, *efficiently evaluate it*

Running example: The simple PROP program:

```
1 (x ∨ y, [x ↦ 0.1, y ↦ 0.3])
```

- Worst-case hardness: computing $\Pr(\varphi_{ex})$ is NP-hard for arbitrary formulae.


We will show this by reduction to the **SAT problem**: the problem of determining whether or not a formula has a model.

Reduction: Let φ be a formula. Assign a probability of 0.5 to every assignment of variables. Then, φ has a model if and only if $\Pr(\varphi) > 0$.

- In fact, computing $\Pr(\varphi_{ex})$ is #P-hard. #P is the class of problems that are polytime reducible to counting the number of solutions to an arbitrary Boolean formula. See Goldreich [2008, Chapter 6] for more discussion on this complexity class, as well as Roth [1996].

Evaluating queries via search

- *Idea:* We want to *search* through the space of models of the formula to potentially avoid exploring all possible worlds. Aim for better average-case/common-case behavior than the naive table enumeration.



- We will describe a relation $(\pi, p) \Downarrow^e \mathbb{R}$:
 - π is an ordered list of propositional variables. We denote list concatenation as $x :: \pi$.
 - p is a PROP program
 - The result \mathbb{R} will be equal to $\llbracket p \rrbracket$
- Define $\varphi[x \mapsto v]$ as the substitution where we replace all instances of x with the value $v \in \{\text{true}, \text{false}\}$ in φ , and “simplify” φ by evaluating connectives until either (1) there are no more true or false constants, or (2) the formula is equal to true or false.
Example: $(x \vee y)[x \mapsto \text{true}] = y$.
- Now, let’s describe our procedure for searching to solve our inference problem. We will define \Downarrow^e inductively. The base-cases will be formulas without any free variables:

$$\begin{array}{ll} \text{(TRUE)} & \text{(FALSE)} \\ (\pi, (\text{true}, w)) \Downarrow^e 1 & (\pi, (\text{false}, w)) \Downarrow^e 0 \end{array}$$

- Then, the inductive step:

$$\frac{\text{(SPLIT)} \quad (\pi, (\varphi[x \mapsto \text{true}], w)) \Downarrow^e v_1 \quad (\pi, (\varphi[x \mapsto \text{false}], w)) \Downarrow^e v_2}{(x :: \pi, (\varphi, w)) \Downarrow^e w(x)v_1 + w(\bar{x})v_2}$$

We are being a bit informal here to save time. If we have time, we can discuss this substitution operation more formally.

- Example derivation tree for our running example for the order $\pi = [x, y]$ and our program $(x \vee y, [x \mapsto 0.1, y \mapsto 0.3])$
- The **cost** of evaluating a formula φ for order π under these semantics is equal to the total number of derivations in the derivation tree.

Theorem 1. Let p be PROP program and π be an ordering on all variables in p , and \Pr be a fully-factorized joint distribution on all variables in φ . If $(\pi, p) \Downarrow v$, then $\llbracket p \rrbracket = v$.

Proof. We will show this by structural induction on p . The base cases are straightforward: If $p = (\text{true}, w)$, then $(\pi, (\text{true}, w)) \Downarrow^e 1$. By definition, $\llbracket (\text{true}, w) \rrbracket = \sum_{I \models \text{true}} w(I) = 1$; similar for the false case.

Induction hypothesis (IH): Let $(x :: \pi, (\varphi[x \mapsto \text{true}], w)) \Downarrow^e v_1$ and $(x :: \pi, (\varphi[x \mapsto \text{false}], w)) \Downarrow^e v_2$, and $(\pi, (\varphi, w)) \Downarrow^e v$. Assume by induction that $v_1 = \llbracket (\varphi[x \mapsto \text{false}], w) \rrbracket$ and $v_2 = \llbracket (\varphi[x \mapsto \text{true}], w) \rrbracket$. Then,

$$\begin{aligned} \llbracket (\varphi, w) \rrbracket &= \sum_{I \in \llbracket \varphi \rrbracket} w(I) \\ &= \underbrace{\sum_{\{I \in \llbracket \varphi \rrbracket, I(x) = \text{true} \}} w(I)}_{(A)} + \underbrace{\sum_{\{I \in \llbracket \varphi \rrbracket, I(x) = \text{false} \}} w(I)}_{(B)} \end{aligned}$$

Example: consider the formula $x \vee y$. Then:

$$\llbracket (x \vee y, w) \rrbracket = \underbrace{w(xy) + w(x\bar{y})}_{(A)} + \underbrace{w(\bar{x}y)}_{(B)}$$

Since $v = w(x)v_1 + w(\bar{x})v_2 = w(x)\llbracket (\varphi[x \mapsto \text{true}], w) \rrbracket + w(\bar{x})\llbracket (\varphi[x \mapsto \text{false}], w) \rrbracket$ by IH, we are done if we can show that $(A) = w(x)\llbracket (\varphi[x \mapsto \text{true}], w) \rrbracket$ and $(B) = w(\bar{x})\llbracket (\varphi[x \mapsto \text{false}], w) \rrbracket$.

To show this observe that after substituting, x is no longer in $\varphi[x \mapsto v]$. Example:

$$\begin{aligned} \llbracket (x \vee y)[x \mapsto \text{false}] \rrbracket &= \llbracket y \rrbracket = w(xy) + w(\bar{x}y) \\ \llbracket (x \vee y)[x \mapsto \text{true}] \rrbracket &= \llbracket \text{true} \rrbracket \end{aligned}$$

Clearly there are permitted models after we restrict in this way. But, due to the factorization of w , we can prove that:

$$\llbracket (\varphi[x \mapsto v], w) \rrbracket = w(x)\llbracket (\varphi, w) \rrbracket.$$

□

- What is the cost of enumerating the formula $x \vee y \vee z \vee w$ with these semantics? It will be linear-time for any order; a big improvement over the exhaustive enumeration!
- However, what about a formula like $(a \wedge b) \vee (c \wedge d)$?

3 Memoization

- This process is sometimes also called *variable elimination*
- Add a context to the reduction relation $\rho : \varphi \rightarrow [0, 1]$ that maps formulae to their probability of being true
- Relation now has the shape:

$$(\pi, \rho, \mathfrak{p}) \Downarrow^m (\mathbb{R}, \rho')$$

References

- Oded Goldreich. Computational complexity: a conceptual perspective. *ACM Sigact News*, 39(3):35–39, 2008.
- Benjamin C Pierce. *Types and programming languages*. MIT press, 2002.
- Dan Roth. On the hardness of approximate reasoning. *Artificial Intelligence*, 82(1-2):273–302, 1996.